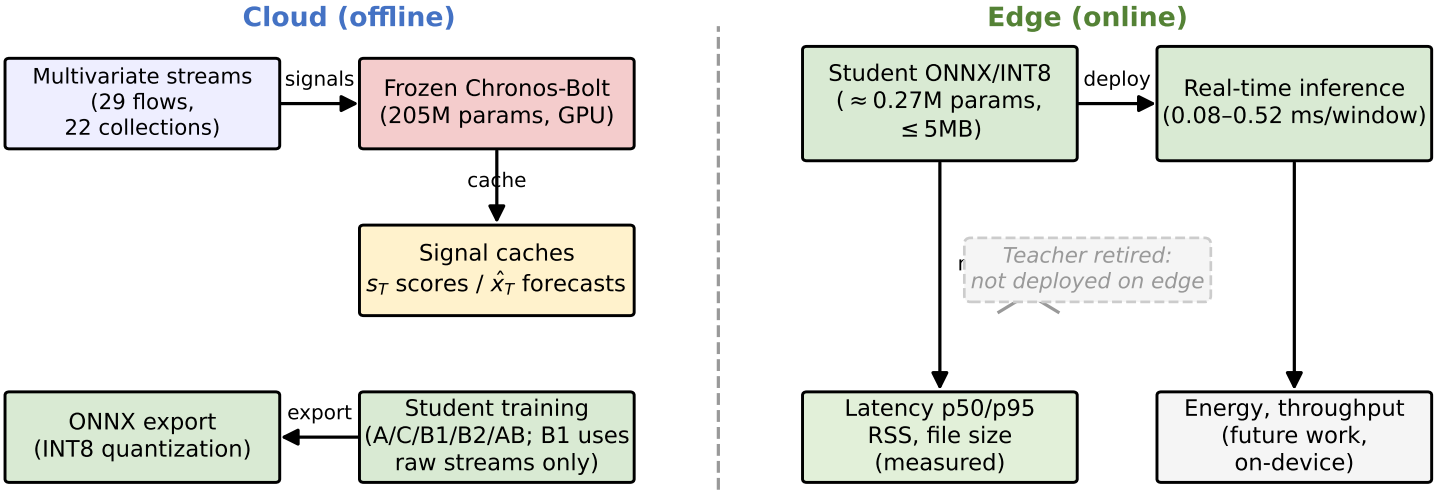
Cloud (offline)
Multivariate streams (29 flows, 22 collections)
signals
Frozen Chronos-Bolt (205M params, GPU)
cache
Signal caches s_T scores / $\hat{x}_T$ forecasts
ONNX export (INT8 quantization)
export
Student training (A/C/B1/B2/AB; B1 uses raw streams only)
Edge (online)
Student ONNX/INT8 (≈ 0.27M params, ≤ 5MB)
deploy
Real-time inference (0.08–0.52 ms/window)
Teacher retired: not deployed on edge
Latency p50/p95 RSS, file size (measured)
Energy, throughput (future work, on-device)